



Deep-learning: music genre recognition

19th October 2018

Gregoire Virepinte, Margaux Wehr, Zhijian Xu, Yuhua Qiao, Ran Huo

Aim – Learnings & Challenges



Classify mp3 tracks into their respective genres

CNN will learn filters that extract features in the frequency and time domain

Learnings

- Converting audio files into visual input
- Pytorch

Challenges

- Large amount of data pre-processing
- Dealing with heavy input size (mp3 format)
- Boundaries between genres can be fuzzy


Overall methodology

Main challenge: transforming audio into usable input for CNN

Data collection



Spectrogram




Numpy array conversion

Slice

Classify using CNN

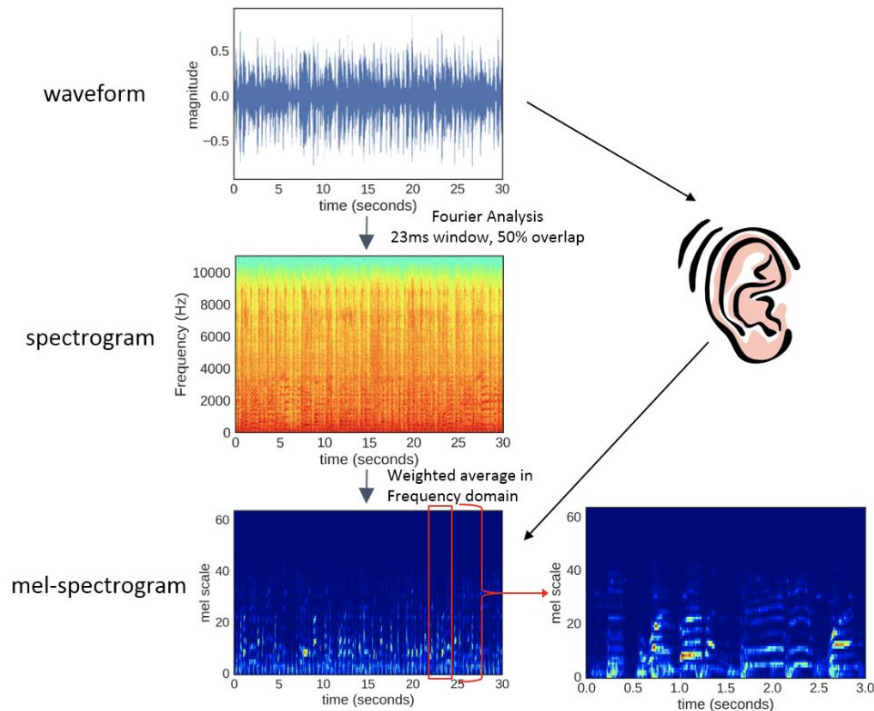
Vote



Answer: Transform audio signal into visual output

Converting audio to image recognition

- A spectrogram is a space, frequency and amplitude representation. It's commonly used to visualize these three attributes of a signal
- **'Divide and conquer'**: split spectrogram into 3s segments
- We use **mel-spectrogram** as the input to the **CNN**





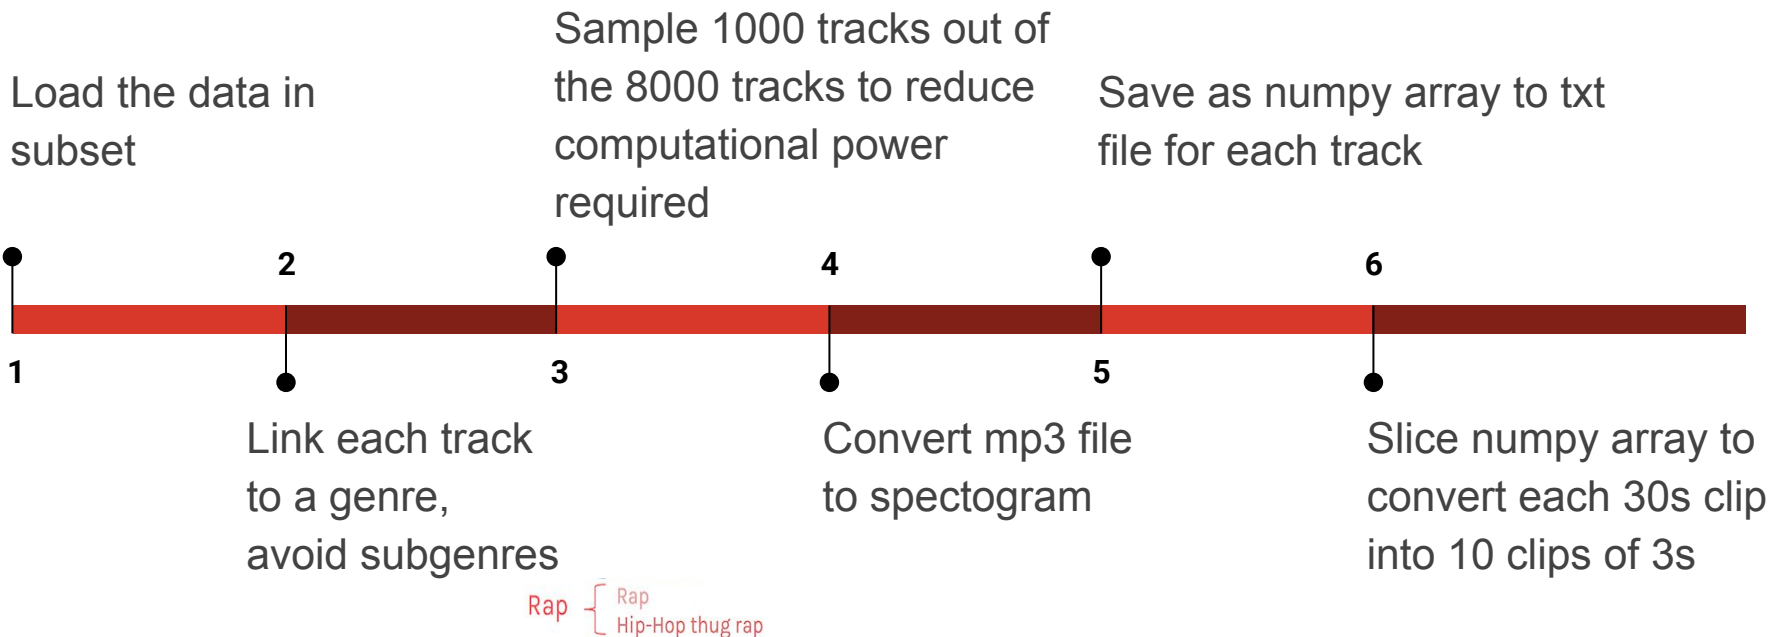
Data collection

- 106,574 of Creative Commons-licenced tracks from 16,341 artists and 14,854 album
- Dataset: **8,000 30-second clips from Free Music Archive**. The clips belong to 8 top genres, balanced with 1,000 clips per genre.
- Our goal in this notebook will be to **classify the genre of an unknown track among the 8 genres that are present in our dataset**.





Data pre-processing





Deep-dive: finding the right genre

- Some songs have **several genres** associated to them
- In the end, the goal of this section was to only keep the "**parent**" genre, which is supposed to be unique, for every track we have in our dataset

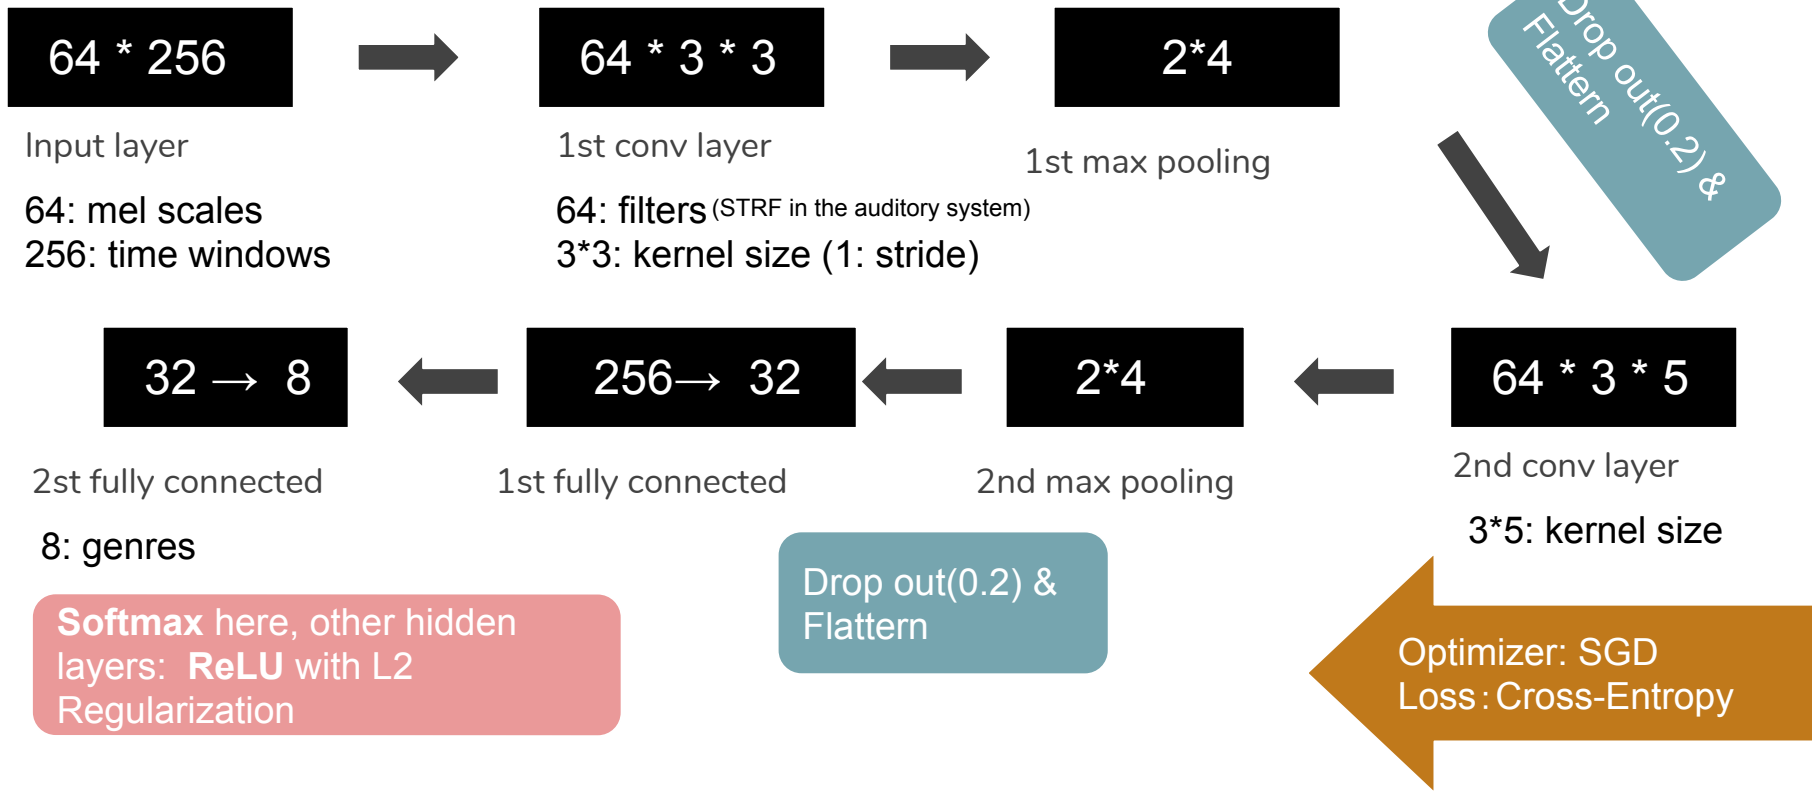
| track_id | genres |
|----------|--------------|
| 2 | Hip-Hop |
| 5 | Hip-Hop |
| 10 | Pop |
| 140 | Folk |
| 141 | Folk |
| 148 | Experimental |
| 182 | Rock |
| 190 | Folk |
| 193 | Folk |
| 194 | Folk |



```
def parent_genre(genre_list):  
    """  
    This is the function we apply on the `genres` column.  
    We replace every genre by its parent genre (if it exists, otherwise we change nothing).  
  
    Afterwards, for a given row, if all parent genres are the same, then we remove  
    the list and keep the first element only.  
    """  
  
    genre_list = genre_list[1:-1].split(',')  
    genre_list = [int(genres.loc[int(x)].top_level) if int(x) in genres.index else int(x) for x in genre_list]  
  
    if len(set(genre_list)) == 1:  
        return genre_list[0]  
  
    else:  
        return genre_list  
  
reduced_tracks.genres = reduced_tracks.genres.apply(parent_genre, 1)
```



CNN: Network Architecture







CNN: Train and Prediction

We set training set, validation set, test set = 7:2:1

Output for the time being:





Improvement

1

Use more data

Run the CNN on the whole dataset, instead of the reduced one

2

Add validation set

Repeat the procedure until classification accuracy on the cross-validation data set doesn't improve anymore

3

Combine the prediction

Combine 10 predictions of 10 3-sec segments of one song
→ probability



Deep learning vs traditional methods

CNN

- Widely used for image classification
- Non-parametric
- No need to input features
- Hard to interpret the model

Random Forest

- "Worry-free" approach, no real hyperparameters to tune
- Non-parametric
- Limit instability of single trees by averaging predictions, limit overfitting
- View feature importance

| | |
|--------------------|-----|
| Train set accuracy | 98% |
| Test set accuracy | 51% |